

Which of the following objects can be shared through secure data sharing?

- A. Masking policy
- B. Stored procedure
- C. Task
- D. External table

Answer: D

Explanation:

The following Snowflake database objects can be shared: Tables External tables Secure views Secure materialized views Secure UDFs

Which formats does Snowflake store unstructured data in? (Choose two.)

- A. GeoJSON
- B. Array
- C. XML
- D. Object
- E. BLOB

Answer: BD

Explanation:

- B. Array
- D. Object

A user is preparing to load data from an external stage.

Which practice will provide the MOST efficient loading performance?

- A. Organize files into logical paths
- B. Store the files on the external stage to ensure caching is maintained
- C. Use pattern matching for regular expression execution
- D. Load the data in one large file

Answer: A

Explanation:

Both internal (i.e. Snowflake) and external (Amazon S3, Google Cloud Storage, or Microsoft Azure) stage references can include a path (or prefix in AWS terminology). When staging regular data sets, we recommend partitioning the data into logical paths that include identifying details such as geographical location or other

What effect does WAIT_FOR_COMPLETION = TRUE have when running an ALTER WAREHOUSE command and changing the warehouse size?

- A. The warehouse size does not change until all queries currently running in the warehouse have completed.
- B. The warehouse size does not change until all queries currently in the warehouse queue have completed.
- C. The warehouse size does not change until the warehouse is suspended and restarted.
- D. It does not return from the command until the warehouse has finished changing its size.

Answer: D

Explanation:

WAIT_FOR_COMPLETION = FALSE | TRUE When resizing a warehouse, you can use this parameter to block the return of the ALTER WAREHOUSE command until the resize has finished provisioning all its compute

Which of the following can be used when unloading data from Snowflake? (Choose two.)

- A. When unloading semi-structured data, it is recommended that the STRIP_OUTER_ARRAY option be used.
- B. Use the ENCODING file format option to change the encoding from the default UTF-8.
- C. The OBJECT_CONSTRUCT function can be used to convert relational data to semi-structured data.
- D. By using the SINGLE = TRUE parameter, a single file up to 5 GB in size can be exported to the storage layer.
- E. Use the PARSE_JSON function to ensure structured data will be unloaded into the VARIANT data type.

Answer: CD

Explanation:

Specifically mentioned in Snowflake Documentation: "Output files are always encoded using UTF-8, regardless of the file format; no other character sets are supported."

What data is stored in the Snowflake storage layer? (Choose two.)

- A. Snowflake parameters
- B. Micro-partitions
- C. Query history
- D. Persisted query results
- E. Standard and secure view results

Answer: BC

Explanation:

B. Micro-partitions
C. Query history (also stored in the data layer of SNOWFLAKE.USA)

A data provider wants to share data with a consumer who does not have a Snowflake account. The provider creates a reader account for the consumer following these steps:

1. Created a user called "CONSUMER"
2. Created a database to hold the share and an extra-small warehouse to query the data
3. Granted the role PUBLIC the following privileges: Usage on the warehouse, database, and schema, and SELECT on all the objects in the share

Based on this configuration what is true of the reader account?

- A. The reader account will automatically use the Standard edition of Snowflake.
- B. The reader account compute will be billed to the provider account.
- C. The reader account can clone data the provider has shared, but cannot re-share it.
- D. The reader account can create a copy of the shared data using CREATE TABLE AS...

Which of the following activities consume virtual warehouse credits in the Snowflake environment? (Choose two.)

- A. Caching query results
- B. Running EXPLAIN and SHOW commands
- C. Cloning a database
- D. Running a custom query
- E. Running COPY commands

Answer: DE

Explanation:

A warehouse provides the required resources, such as CPU, memory, and temporary storage, to perform the

When loading data into Snowflake, the COPY command supports which of the following?

- A. Joins
- B. Filters
- C. Column reordering
- D. Aggregates

Answer: C

Explanation:

Reference:

<https://docs.snowflake.com/en/user-guide/data-load-transform.html#reorder-csv-columns-during-a-load>

What is cached during a query on a virtual warehouse?

- A. All columns in a micro-partition
- B. Any columns accessed during the query
- C. The columns in the result set of the query
- D. All rows accessed during the query

Answer: B

Explanation:

B“As Snowflake is a columnar data warehouse, it automatically returns the columns needed rather than the entire row to further help maximize query performance”

<https://community.snowflake.com/s/article/Caching-in-Snowflake-Data-Warehouse>

What is the default character set used when loading CSV files into Snowflake?

- A. UTF-8
- B. UTF-16
-  C. ISO 8859-1

Question: 355

Which of the following describes external functions in Snowflake?

- A. They are a type of User-defined Function (UDF).
- B. They contain their own SQL code.
- C. They call code that is stored inside of Snowflake.
- D. They can return multiple rows for each row received.

Answer: A

Explanation:

A correctAn external function is a type of UDF. Unlike other UDFs, an external function does not contain its own code; instead, the external function calls code that is stored and executed out

Question: 351**SnoPro**

Which of the following activities consume virtual warehouse credits in the Snowflake environment? (Choose two.)

- A. Caching query results
- B. Running EXPLAIN and SHOW commands
- C. Cloning a database
- D. Running a custom query
- E. Running COPY commands

Answer: DE**Explanation:**

A warehouse provides the required resources, such as CPU, memory, and temporary storage, to perform the

Question: 356**SnoPro**

Which of the following are valid methods for authenticating users for access into Snowflake? (Choose three.)

- A. SCIM
- B. Federated authentication
- C. TLS 1.2
- D. Key-pair authentication
- E. OAuth
- F. OCSP authentication

Answer: BDE**Explanation:**

BDE

<https://docs.snowflake.com/en/user-guide/authentication.html>

A user has a standard multi-cluster warehouse auto-scaling policy in place.

Which condition will trigger a cluster to shut-down?

[Download to res](#)

- A. When after 2-3 consecutive checks the system determines that the load on the most-loaded cluster could be redistributed.
- B. When after 5-6 consecutive checks the system determines that the load on the most-loaded cluster could be redistributed.
- C. When after 5-6 consecutive checks the system determines that the load on the least-loaded cluster could be redistributed.
- D. When after 2-3 consecutive checks the system determines that the load on the least-loaded cluster could be redistributed.

Question: 358

SnoPro

What is the minimum Snowflake edition needed for database failover and fail-back between Snowflake accounts for business continuity and disaster recovery?

- A. Standard
- B. Enterprise
- C. Business Critical
- D. Virtual Private Snowflake

Answer: C

Explanation:

Business Critical

Requires Business Critical (or higher).

<https://docs.snowflake.com/en/user-guide/database-failover-config.html#failing-over-databases-across-multiple-accounts>

Question: 359**SnoPro**

How would a user execute a series of SQL statements using a task?

- A. Include the SQL statements in the body of the task `CREATE TASK mytask .. AS INSERT INTO target1 SELECT .. FROM stream_s1 WHERE .. INSERT INTO target2 SELECT .. FROM stream_s1 WHERE ..`
- B. A stored procedure can have only one DML statement per stored procedure invocation and therefore the user should sequence stored procedure calls in the task definition `CREATE TASK mytaskAS call stored_proc1(); call stored_proc2();`
- C. Use a stored procedure executing multiple SQL statements and invoke the stored procedure from the task. `CREATE TASK mytaskAS call stored_proc_multiple_statements_inside();`
- D. Create a task for each SQL statement (e.g. resulting in task1, task2, etc.) and string the series of SQL statements by having a control task calling task1, task2, etc. sequentially.

Question: 360

How many resource monitors can be assigned at the account level?

- A.1
- B.2
- C.3
- D.4

Answer: A

Explanation:

A single monitor can be set at the account level to control credit usage for all warehouses in

Question: 361

Question: 362

Data storage for individual tables can be monitored using which commands and/or objects? (Choose two.)

- A.SHOW STORAGE BY TABLE;
- B.SHOW TABLES;
- C.Information Schema -> TABLE_HISTORY
- D.Information Schema -> TABLE_FUNCTION
- E.Information Schema -> TABLE_STORAGE_METRICS

Answer: BE

Explanation:

B. SHOW TABLES. E. TABLE_STORAGE_METRICSThese two options will show bytes stored.

<https://docs.snowflake.com/en/sql-reference/sql/show-tables.html>

https://docs.snowflake.com/en/sql-reference/info-schema/table_storage_metrics.html

Question: 362**SnoPro**

How would a user run a multi-cluster warehouse in maximized mode?

- A. Configure the maximum clusters setting to "Maximum."
- B. Turn on the additional clusters manually after starting the warehouse.

[Download to](#)

- C. Set the minimum Clusters and maximum Clusters settings to the same value.
- D. Set the minimum clusters and maximum clusters settings to different values.

Answer: C**Explanation:**

It is C. If min=max, there is no room for increasing any clusters and min and max would be same. Hence same

Question: 363**SnoPro**

What internal stages are available in Snowflake? (Choose three.)

- A. Schema stage
- B. Named stage
- C. User stage
- D. Stream stage
- E. Table stage
- F. Database stage

Answer: BCE**Explanation:**

BCE User stage, Table Stage and Named stage are internal stages of snowflake,

<https://hevodata.com/learn/snowflake-stages/#ist>

**Question: 364****SnoPro**

Which stages are used with the Snowflake PUT command to upload files from a local file system? (Choose three.)

- A. Schema Stage
- B. User Stage**
- C. Database Stage
- D. Table Stage**
- E. External Named Stage
- F. Internal Named Stage**

Answer: BDF**Explanation:**

BDFUploads (i.e. stages) data files from a local directory/folder on a client machine to one of the following Snowflake stages:Named internal stage.Internal stage for a specified table.Internal stage for the current user.

**Question: 590****SnoPro**

Which SQL commands should be used to write a recursive query if the number of levels is unknown? (Choose two.)

- A.CONNECT BY**
- B.LISTAGG
- C.MATCH RECOGNIZE
- D.QUALIFY
- E.WITH**

Answer: AE**Explanation:**

CONNECT BY , the recursive CTE portion of the WITH command , Working with CTEs (Common Table Expressions) , Tabular SQL UDFs (UDTFs)

Question: 691**SnoPro**

What information is stored in the ACCESS_HISTORY view?

- A. History of the files that have been loaded into Snowflake
- B. Names and owners of the roles that are currently enabled in the session
- C. Query details such as the objects included and the user who executed the query**
- D. Details around the privileges that have been granted for all objects in an account

Answer: C**Explanation:**

https://docs.snowflake.com/en/sql-reference/account-usage/access_history#columns

Question: 592


SnowPro Core: Actual Exam Q&A | CLEARCATNET

What privilege does a user need in order to receive or request data from the Snowflake Marketplace?

- A.CREATE DATA EXCHANGE LISTING
- B.CREATE SHARE
- ☒ C.IMPORT SHARE



[Download to read ad-free](#)



D.IMPORTED PRIVILEGES

Question: 593

SnoPro

Which Snowflake database object can be shared with other accounts?

- A.Tasks
- B.Pipes
- C.Secure User-Defined Functions (UDFs)**
- D.Stored Procedures

Answer: C

Explanation:

Reference:

<https://docs.snowflake.com/en/user-guide/data-sharing-intro> You can share the following Snowflake database objects: Tables External tables Secure views Secure materialized views Secure UDFs

Which identity providers are valid type values for federated authentication on the SAML_IDENTITY_PROVIDER parameter? (Choose two.)

- A. Identity Access Management (IAM)
- B. Microsoft Active Directory Federation Services (AD FS)**
- C. OAuth
- D. Okta**
- E. PingFederate

Answer: BD

Explanation:

B,D. But this parameter is deprecated

<https://docs.snowflake.com/ko/sql-reference/parameters#label-saml-identity-provider>

A Snowflake user wants to share data using my_share with account xy12345.

Which command should be used?

- A. grant usage on share my_share to account xy12345;
- B. grant select on share my_share to account xy12345;
- C. alter share my_share add accounts = xy12345;
- D. alter account xy12345 add share my_share;

Answer: C

Explanation:

```
ALTER SHARE [ IF EXISTS ] <name> ADD | REMOVE ACCOUNTS = <consumer_account> [ ,  
<consumer_account> , ... ][ SHARE_RESTRICTIONS = TRUE | FALSE ]ALTER SHARE [ IF  
EXISTS ] <name> SET [ ACCOUNTS = <consumer_account> [ , <consumer_account> ... ] ]
```

What role is required to use Partner Connect?

- A.ACCOUNTADMIN
- B.ORGADMIN
- C.SECURITYADMIN
- D.SYSADMIN

Answer: A

Explanation:

Partner Connect is limited to account administrators (i.e. users with the ACCOUNTADMIN role) who have a verified email address in Snowflake: To use Partner Connect, you must switch to the ACCOUNTADMIN role or contact someone in your organization who has the role.

How can a Snowflake user configure a virtual warehouse to support over 100 users if their company has Enterprise Edition?

- A. Add additional warehouses and configure them as a cluster.
- B. Set the auto-scale to 100.
- C. Use a multi-cluster warehouse.
- D. Use a larger warehouse.

Answer: C

Explanation:

Use Multi-Cluster Warehouses for concurrency

How is table data compressed in Snowflake?

- A. Each column is compressed as it is stored in a micro-partition.
- B. Each micro-partition is compressed as it is written into cloud storage using GZIP.
- C. The micro-partitions are stored in compressed cloud storage and the cloud storage handles compression.
- D. The text data in a micro-partition is compressed with GZIP but other types are not compressed.

Answer: A

Explanation:

<https://docs.snowflake.com/en/user-guide/tables-clustering-micropartitions>

COLUMNAR FORMAT When data is loaded into Snowflake, Snowflake reorganizes that data into its internal optimized, compressed, columnar format. Snowflake stores this optimized data in cloud storage. Snowflake manages all aspects of how this data is stored – the organization, file size, structure, compression, metadata, statistics, and other aspects of data storage are handled by Snowflake. The data objects stored by Snowflake are not directly visible nor accessible by customers; they are only accessible through SQL query operations run using Snowflake.

Question: 599

SnoPro

What will be the output of the below query against the table name gold_data?

```
select * from gold_data tablesample (100);
```

- A. It will return an empty sample.
- B. It will return a random 100 rows.
- C. It will return an entire table.
- D. It will produce an error message.

Answer: C

Explanation:

<https://docs.snowflake.com/en/sql-reference/constructs/sample>

from doc:Return an entire table, including all rows in the table:SELECT * FROM testtable TABLESAMPLE (100);

Question: 600

SnoPro

A Snowflake query took 40 minutes to run. The results indicate that 'Bytes spilled to local storage' was a large number.

What is the issue and how can it be resolved?

- A. The warehouse is too large. Decrease the size of the warehouse to reduce the spillage.
- B. The warehouse is too small. Increase the size of the warehouse to reduce the spillage.
- C. The Snowflake console has timed-out. Contact Snowflake Support.



Question: 601

SnoPro

What is the MOST efficient way to load streaming data into Snowflake?

- A. Use the COPY command.
- B. Use Snowpipe.
- C. Use the Load Data Wizard.
- D. Use tasks and streams.

Answer: B

Explanation:

The MOST efficient way to load streaming data into Snowflake is to use Snowpipe (option B). Snowpipe is a native, automated data ingestion service within Snowflake that can continuously load data from various sources in near real-time.

Which COPY INTO statement accurately describes how to unload data from a Snowflake table?

- A. The default value for the SINGLE option is set to TRUE.
- B. By default, COPY INTO statements do not separate table data into a set of output files.
- C. The OBJECT_CONSTRUCT function can be combined with the COPY command to convert the rows in a relational table to a single VARIANT column.
- D. If the COMPRESSION option is set to TRUE, a file's name can be specified with the appropriate file extension so that the output file can be compressed.

Answer: C

Explanation:


"You can use the OBJECT_CONSTRUCT function combined with the COPY command to convert the rows in a relational table to a single VARIANT column and unload the rows into a file.

<https://docs.snowflake.com/en/user-guide/data-unload-considerations>

Question: 603

Sr

What command is used to download data from a Snowflake stage?



- A.PUT
- B.INSERT
- C.GET
- D.COPY

Question: 604

Snor

By default, which role has privileges to create tables and views in an account?

- A.PUBLIC
- B.SECURITYADMIN
- C.SYSADMIN
- D.USERADMIN

Answer: C

Explanation:

The SYSADMIN role is a system-defined role that has privileges to create warehouses, databases, and database objects in an account and grant those privileges to other roles.

<https://docs.snowflake.com/en/user-guide/security-access-control-configure#creating-a-role-hierarchy>

What does Snowflake recommend as a best practice for using secure views?

- A. Use sequence-generated values.
- B. Programmatically reveal the identifiers.
- C. Use secure views solely for query convenience.
- D. Do not expose the sequence-generated column(s).

Answer: D

Explanation:

Sequence-generated Columns Do not expose the sequence-generated column as part of the view.

Reference:

<https://docs.snowflake.com/en/user-guide/views-secure>

Question: 606

SnowPro Core: Actual Exam Q&A | CLEARCATNET

What is the Fail-safe period for a transient table in the Snowflake Enterprise edition and higher?

- A. 0 days
- B. 1 day
- C. 7 days
- D. 14 days

Answer: A

Explanation:

Temporary and Transient Tables To help manage the storage costs associated with Time Travel and Fail-safe, Snowflake provides two table types, temporary and transient, which do not incur the same fees as standard (i.e. permanent) tables: Transient tables can have a Time Travel retention period of either 0 or 1 day. Temporary tables can also have a Time Travel retention period of 0 or 1 day; however, this retention period ends as soon

How does a Snowflake user enable Multi-Factor Authentication (MFA)?

- A. The user must enroll themselves through the web interface.
- B. The user must submit their encrypted private key to Snowflake.
- C. The user must sign up with Duo Mobile for approval to use the service.
- D. The user must configure Snowflake to use Single Sign-On (SSO).

Answer: A

Explanation:

A is correct.

<https://docs.snowflake.com/en/user-guide/security-mfa>

<https://docs.snowflake.com/en/user-guide/security-mfa#:~:text=Any%20Snowflake%20user%20can%20self%2Denroll%20in%20MFA%20through%20the%20web%20interface>

Question: 608

SnoPro

What allows a user to limit the number of credits consumed within a Snowflake account?

- A. Tracking account usage
- B. Creating resource monitors
- C. Automatic virtual warehouse scaling
- D. Automatic clustering

Question: 609

SnoPro

Which statement accurately describes Snowflake's architecture?

- A. It uses a local data repository for all compute nodes in the platform.
- B. It is a blend of shared-disk and shared-everything database architectures.
- C. It is a hybrid of traditional shared-disk and shared-nothing database architectures.
- D. It reorganizes loaded data into internal optimized, compressed, and row-based format.

Answer: C

Explanation:

Snowflake architecture base question

**Question: 610****SnoPro**

Which Snowflake SQL command is used to get a subset of rows randomly from a table?

- A.GENERATOR
- B.LATERAL
- C.PIVOT
- D.SAMPLE

Answer: D**Explanation:**

Reference:

<https://docs.snowflake.com/en/sql-reference/constructs/sample>

Question: 611**SnoPro**

Which statement accurately describes how a virtual warehouse functions?

- A. Increasing the size of a virtual warehouse will always improve data loading performance.
- B. Each virtual warehouse is an independent compute cluster that shares compute resources with other warehouses.
- C. Each virtual warehouse is a compute cluster composed of multiple compute nodes allocated by Snowflake from a cloud provider.
- D. All virtual warehouses share the same compute resources so performance degradation of one warehouse can significantly affect all the other warehouses.

<https://quizlet.com/380510774/snowflake-certification-flash-cards/>

